

Introduction to parallel computing with IPython

- Credits for a [prior version of this document go to Skipper Seabold](#)
- Most of this presentation is taken from the [IPython parallel computing documentation](#)
- And from talks given over the years by the core development team of [@minrk](#), [@ellisonbg](#), and [@fperez_org](#), among many others
- IPython is well [documented](#), including [video tutorials](#)
- There is a great support network for IPython on [stackoverflow](#) and on their [mailing list](#)
- This talk is created using the [IPython Notebook](#), which also support parallelism


```
In [1]: %pylab inline
import numpy as np
import matplotlib.pyplot as plt
```

```
Welcome to pylab, a matplotlib-based Python environment [backend:
module://IPython.kernel.zmq.pylab.backend_inline].
For more information, type 'help(pylab)'.
```

Architecture

```
In [2]: from IPython.core.display import Image
        Image("parallel_architecture400.png")
```

Out[2]:



Three Core Parts

I. The **Client**

- This is what you use to run your parallel computations
- You will interact with a **View** on the client
- The type of View depends on the execution model you are using

II. The **Controller**

- The IPython controller consists of 1) the **Hub** and 2) the **schedulers**
- The **Hub** is the central process that monitors everything
- The **schedulers** take care of getting your code where it should go
- The controller is the go between from the Client to the Engines

III. The **Engine**

- An IPython "kernel" where the code is executed
- Listens for instructions over a network connection

IPython client and views

- The **Client** object connects to the cluster
- For each execution model there is a corresponding **View**

- For example, there are two basic views:
 - The `DirectView` class for explicitly running code on a particular engine(s)
 - The `LoadBalancedView` class for running your code on the 'best' engine(s)
- You can use as many views as you like, many at the same time
- You can read much more about the details of IPython parallel and views in the [documentation](#)

Getting Started

- Take advantage of multiple the processors on your local machine
- Say you have 4 processors
- How many processors do I have available?

```
In [3]: from multiprocessing import cpu_count
        print cpu_count()
```

4

- Start a controller and 4 engines with **ipcluster**
- At the command line type

```
ipcluster start -n 4
```

- Or, in the notebook at the dashboard

Did it work?

```
In [4]: from IPython import parallel

        rc = parallel.Client()
        rc.block = True
```

```
In [5]: def power(a, b):
        return a**b
```

- Create a direct view of kernel 0
- Direct views support slicing

```
In [6]: dv = rc[0]
        dv
```

```
Out[6]: <DirectView 0>
```

```
In [7]: dv.apply(power, 2, 10)
```

```
Out[7]: 1024
```

Recall that slice notation allows you leave out start and stop steps

```
In [8]: x = [1, 2, 3, 4]
x
```

```
Out[8]: [1, 2, 3, 4]
```

```
In [9]: x[:]
```

```
Out[9]: [1, 2, 3, 4]
```

Use this to send code to all the engines

```
In [10]: rc[:].apply_sync(power, 2, 10)
```

```
Out[10]: [1024, 1024, 1024,
          1024]
```

Python's built-in map function allows you to call a sequence a function over a sequences of arguments

```
In [11]: map(power, [2]*10, range(10))
```

```
Out[11]: [1, 2, 4, 8, 16, 32, 64, 128, 256,
          512]
```

In parallel, you use view.map

```
In [12]: view = rc.load_balanced_view()
view.map(power, [2]*10, range(10))
```

```
Out[12]: [1, 2, 4, 8, 16, 32, 64, 128, 256,
          512]
```

The Direct Interface

- The direct interface lets the user interact explicitly with each engine
- First, create a direct view

```
In [13]: from IPython import parallel

rc = parallel.Client()
```

- Above we saw the use of map and apply in parallel
- You may have also noticed this bit of code

```
In [14]: rc.block = True
```

- In blocking mode, whenever you call execute code on the engines the controller waits until this code is done executing
- Non-blocking mode is the default

- Get access to all the engines

```
In [15]: dview = rc[:]
```

- You can block on a call-by-call basis as well, by using `apply_sync` for synchronous execution

```
In [16]: dview.block = False

dview["a"] = 5 # shorthand for push
dview["b"] = 7

dview.apply_sync(lambda x: a + b + x, 27)
```

```
Out[16]: [39, 39, 39,
          39]
```

- There is also `apply_async`
- Above you'll notice that the assignments defined these variables on the engines in a dictionary-like manner
- This is shorthand for pushing python objects to the engines
- DirectViews provide dictionary-like access by key or by using `get` and `update` like built-in dicts
- This can also be done explicitly with `push`
- `push` takes a dictionary

```
In [17]: dview.push(dict(msg="Hi, there"), block=True)
```

```
Out[17]: [None, None, None,
          None]
```

```
In [18]: dview.block = True
```

- Python commands can be executed as strings on specific engines

```
In [19]: dview.execute("x = msg")
```

```
Out[19]: <AsyncResult:
         finished>
```

Or using the interactive `%px` magic, short for "parallel execute"

```
In [20]: %px y = msg + ' you'
```

```
In [21]: print dview["x"] # shorthand for pull
         print dview["y"]

['Hi, there', 'Hi, there', 'Hi, there', 'Hi, there']
['Hi, there you', 'Hi, there you', 'Hi, there you', 'Hi, there you']
```

- You can also pull results back from the engine

```
In [22]: rc[:,2].execute("c = a + b")
         rc[1:,2].execute("c = a - b")
```

```
Out[22]: <AsyncResult:
         finished>
```

```
In [23]: dview.pull("c")
```

```
Out[23]: [12, -2, 12,
         -2]
```

- If we were working in non-blocking mode, we would get an [AsyncResult object](#) back immediately

```
In [24]: def wait(t):
         import time
         tic = time.time()
         time.sleep(t)
         return time.time() - tic
```

```
In [25]: ar = dview.apply_async(wait, 2)
```

```
In [26]: type(ar)
```

```
Out[26]: IPython.parallel.client.asyncresult.AsyncResult
```

- We use its `get` method to get the result
- Calling `get` blocks

```
In [27]: ar.get()
```

```
Out[27]: [2.0022220611572266,
          2.0020439624786377,
          2.0007009506225586,
          2.0022499561309814]
```

- If we weren't quite so patient, we could ask if our tasks are done by using the ready method

```
In [28]: ar = dview.apply_async(wait, 15)
         print ar.ready()
False
```

- Or we can ask for the result, waiting a maximum of, say, 5 seconds

```
In [29]: ar.get(5)
```

```
-----
TimeoutError                                Traceback (most recent call last)
<ipython-input-29-8ed98da2cafb> in <module>()
----> 1 ar.get(5)

/home/fperez/usr/lib/python2.7/site-packages/IPython/parallel/client
/asyncresult.pyc in get(self, timeout)
    126         raise self._exception
    127     else:
--> 128         raise error.TimeoutError("Result not ready.")
    129
    130     def _check_ready(self):

TimeoutError: Result not ready.
```

- Often, we can't go on until some results are done
- For this, we can use the wait method
- wait can take an iterable of AsyncResults

```
In [30]: result_list = [dview.apply_async(wait, 1) for i in range(8)]
```

```
In [31]: dview.wait(result_list)
```

```
Out[31]: True
```

```
In [32]: result_list[0].get()
```

```
Out[32]: [1.0011539459228516,  
          1.0010430812835693,  
          1.0004119873046875,  
          1.0010390281677246]
```

Connecting directly to your engines

Every IPython engine can be turned into a kernel in one command:

```
In [33]: %%px  
         from IPython.parallel import bind_kernel  
         bind_kernel()
```

Let's send the engine IDs to all of them

```
In [34]: dview.scatter('eid', rc.ids, flatten=True)
```

And open a console directly on one of them:

```
In [35]: rc[0].execute("%qtconsole")
```

```
Out[35]: <AsyncResult:  
         finished>
```

Scatter and Gather

- You can use `scatter` to partition an iterable across engines
- `gather` pulls the results back
- You can use this to do parallel list comprehensions as below
- Sometimes this is more convenient than `map`

```
In [36]: dview.scatter('x', range(64))
```

```
%%px y = [i**10 for i in x]
```

```
y = dview.gather('y')
```

```
print y[:10]
```

```
[0, 1, 1024, 59049, 1048576, 9765625, 60466176, 282475249, 1073741824, 3486784401]
```


- The % indicates that we are using an [IPython 'magic' function](#)
- The available parallel magics are listed in the [documentation](#)

The Task Interface

- [The Task interface](#) allows you to use your engines as a system of workers
- You no longer have direct access to the individual engines
- If your tasks are easily segmented into pieces that do not depend on each other, the Task Interface may be ideal
- *However*, you can specify complex dependencies to describe task execution order
- You can use many standard scheduling paradigms for how tasks should be run or define your own
- I am not going to discuss the task interface in detail

```
In [37]: rc = parallel.Client()

         lview = rc.load_balanced_view()
```

```
In [38]: lview.block = True

         parallel_result = lview.map(lambda x:x**10, range(32))

         print parallel_result[:10]

[0, 1, 1024, 59049, 1048576, 9765625, 60466176, 282475249, 1073741824, 3486784401]
```

There's much more

- [Parallel function decorators](#)
- [Using MPI for message passing](#) (or pyzmq)
- [Enabling database backends for storing information on running jobs to disk](#)
- [DAGs for tasks](#)